1. Which of our open ML Research Engineer roles are you interested in, in order?

    1. Model Evaluation
    2. General
    3. Applied Research

Model evaluation combines technical and analytical thinking with the high-level solutions needed for the success of WeatherMesh and Windborne as a whole, and that sort of solutions-based technical thinking is the sweet spot for my brain. The General and Applied Research roles are also aligned with ML research I've done in the past and enjoyed!

2. Describe an ML idea you were initially excited about but later decided was wrong or unimportant. What changed your mind?

My friends and I tried to create a novel pipeline for faster image segmentation by assigning a material classification to clusters of pixels, then merging them. After building the classifier and designing a merging pipeline to create image masks, we realized a segmentation pipeline involving material classification needs to be trained end-to-end using those material predictions. We saw our idea through to the end, but the accuracy numbers proved us wrong!

3. Describe a time when you embarked on a sidequest / took initiative outside your core job responsibilities. What value did you add to your organization? What drove you to action?

At my current internship, I was asked to build a database to quantify the impacts of wildfires on water utilities. Realizing that this would be a great resource that few would use, I developed an interactive Streamlit dashboard to visualize data insights and conducted user research without being asked. This created a product usable by teams across my org. I was driven to action because everyone deserves to and can have access to high-quality data-driven wildfire insights.

4. How could better accuracy fail to translate into more useful forecasts for a weather model?

Assuming "accuracy" means single-variable accuracy in weather models and "useful" means providing the most value to the most people, the translation of atmospheric data into weather forecasts which are relevant and communicable is a non-trivial task. Deficits in interpretability, long lag times in creating forecasts from data, and misunderstanding client needs can waste accurate weather model data by failing to anticipate what a user actually wants from the model.

5. Describe a time you changed how you use LLMs in your work — switching models, tools, or workflows. What did you observe that prompted the change?

I've noticed how much I can learn about data tooling processes by reading the reasoning traces in Cursor as the LLM works. My work this summer has led me to some new tools such as

GeoParquet, and I observed that GPT 5.6 Sol reasons about the implementation process in ways that give me insights into the use cases of tools like GeoParquet. Reading through its implementation strategy has taught me more about how to consider the tradeoffs of using different tools in future projects.

6. *(Optional)* Your fav ML Twitter account?

On Twitter he's @erikphoel, but I read his Substack called the Intrinsic Perspective. He's more of a general AI writer than an ML-specific writer, but he's informed the way I think about model tradeoffs and inspired some rabbit holes I've been down recently.